

04 FOS Marketing and Communications Technical Specification

Founder Operating System

Marketing and Communications Operating System -
Complete Technical Specification and Implementation Plan

1. Implementation directive
2. Business objectives
 - 2.1 Primary objective
 - 2.2 Secondary objectives
 - 2.3 Core metrics
3. Scope
 - 3.1 Included channels and assets
 - 3.2 Out of scope during beta
4. Phase placement
5. Architecture
 - 5.1 Required shared services
6. Roles and permissions
 - Founder
 - Marketing Strategy Agent
 - Content Agents
 - Research Agent
 - Engagement Agent
 - Publication Adapter

7. Domain model

- 7.1 AudienceSegment
- 7.2 ContentPillar
- 7.3 Narrative
- 7.4 FounderVoicePolicy
- 7.5 ChannelPolicy
- 7.6 CallToAction
- 7.7 ContentOpportunity
- 7.8 ResearchBrief
- 7.9 SourceRecord
- 7.10 EvidenceMatrix
- 7.11 SourceBrief
- 7.12 Campaign
- 7.13 ContentAsset
- 7.14 EditorialCalendarItem
- 7.15 PublicationRecord
- 7.16 EngagementRecord
- 7.17 ContentPerformance
- 7.18 ContentExperiment

8. Agent specifications

- 8.1 Editorial Strategy Agent
- 8.2 LinkedIn Content Agent
- 8.3 Substack Research and Essay Agent
- 8.4 Campaign Planning Agent
- 8.5 Repurposing Agent
- 8.6 Claims, Consent, Pricing, and CTA Evaluator
- 8.7 Founder Voice Evaluator
- 8.8 Engagement Intelligence Agent

9. Workflows

- 9.1 Content opportunity intake
- 9.2 LinkedIn post workflow
- 9.3 LinkedIn carousel workflow
- 9.4 Substack paper workflow
- 9.5 Beta launch campaign workflow
- 9.6 Release communications workflow

9.7 Engagement workflow

10. APIs

Strategy and configuration

Opportunities and research

Campaigns and assets

Calendar, publication, and performance

11. User interface

11.1 Marketing Command Center

11.2 Editorial Calendar

11.3 LinkedIn Studio

11.4 Substack Paper Workspace

11.5 Campaign Workspace

12. Background jobs

13. Deterministic safeguards

14. Testing specification

14.1 Unit tests

14.2 Integration tests

14.3 Agent contract fixtures

14.4 End-to-end scenarios

E2E-MCOM-001: Substack to enrollment campaign

E2E-MCOM-002: LinkedIn carousel

E2E-MCOM-003: Release communication

E2E-MCOM-004: Consent revocation

15. Metrics and evaluation

Agent-level

Workflow-level

Business-level

16. Feature flags

17. Work packages

WP-MCOM-0.1 - Strategy and configuration foundation

WP-MCOM-0.2 - Claims, consent, pricing, and publication gates

WP-MCOM-1.1 - Beta launch campaign

WP-MCOM-1.2 - LinkedIn operating workflow

- WP-MCOM-1.3 - Substack research and publication workflow
- WP-MCOM-2.1 - Recurring editorial engine
- WP-MCOM-2.2 - Repurposing and engagement intelligence
- WP-MCOM-3.1 - Release, case study, and technical paper workflows
- WP-MCOM-4.1 - Attribution, experiments, and optimization

- 18. Deployment sequence
- 19. Definition of done
- 20. Coding-agent execution instruction

Founder Operating System

Marketing and Communications Operating System - Complete Technical Specification and Implementation Plan

Document control	Value
Document ID	FOS-MCOM
Version	2.0
Status	Implementation specification
Product owner	Founder
Primary audience	Founder, coding agents, marketing operations, product reviewers
Updated	2026-07-13
Dependencies	FOS core memory, evidence, claims, approvals, events, identity, agent runtime

This document makes marketing and communications a first-class Founder Operating System subsystem. It covers LinkedIn posts and carousels, Substack papers and newsletters, website

and landing-page copy, email sequences, webinars, release communications, build logs, engagement intelligence, repurposing, approval, publication, and attribution.

1. Implementation directive

Build a Marketing and Communications Operating System, or MCOM, inside the existing FOS. MCOM must transform approved strategy, product evidence, founder ideas, user questions, beta learning, releases, and market research into coordinated communications that generate qualified demand and reduce founder production time.

MCOM must not be implemented as a generic content generator. It must preserve:

- Audience intent
- Founder point of view
- Evidence provenance
- Product and pricing accuracy
- Customer consent
- Channel-native structure
- Campaign sequence
- Approval history
- Enrollment attribution
- Founder edits and voice learning

During beta, MCOM may create internal and platform-native drafts but may not publish autonomously.

2. Business objectives

2.1 Primary objective

Increase qualified beta applications and enrollments attributable to founder communications while reducing founder hours per published asset.

2.2 Secondary objectives

- Establish a consistent founder editorial presence.
- Publish substantive long-form work, not only short promotional posts.
- Convert one strong source artifact into multiple channel-native assets.
- Ensure every factual claim is supported.
- Turn customer questions and product learning into useful public education.
- Identify which narratives, formats, channels, and calls to action produce qualified demand.
- Preserve an auditable record from source evidence to publication and business outcome.

2.3 Core metrics

- Qualified leads per asset
- Applications per campaign
- Calls booked per asset
- Enrollments and net revenue per campaign
- Content-assisted conversion
- Subscriber growth
- Founder hours per asset
- Draft approval-with-minor-edits rate
- Repurposing yield per source artifact
- Unsupported-claim block rate

- Publication consistency
- Audience-quality score

3. Scope

3.1 Included channels and assets

- LinkedIn text posts
- LinkedIn document carousels
- LinkedIn launch and follow-up sequences
- Substack research papers
- Substack newsletters
- Website pages
- Enrollment landing pages
- Product pages
- Email campaigns and nurture sequences
- Webinar invitations, scripts, and follow-up
- Build logs
- Release notes and public changelogs
- Case studies
- Technical architecture papers
- FAQ and objection content
- Lead magnets and downloadable guides
- Comment and audience-response drafts

3.2 Out of scope during beta


- Autonomous publishing
- Autonomous direct messages
- Automated comment posting
- Unapproved paid-media spending
- Automated pricing changes
- Fabricated customer stories
- Purchased or scraped contact lists

- Impersonation of the founder
- Crisis communications without founder control
- Competitor allegations without evidence and legal review where required

4. Phase placement

Phase	MCOM deliverable
Phase 0	Audience, voice, brand, content pillars, channels, CTAs, claims, campaign and attribution foundations
Phase 1	Beta launch campaign, LinkedIn sequence, Substack cornerstone paper, landing-page and webinar communications
Phase 2	Recurring founder editorial engine, content calendar, repurposing, engagement intelligence
Phase 3	Release communications, beta-learning reports, case-study evidence, technical papers
Phase 4	Scaled campaign orchestration, experimentation, attribution, optimization
Phase 5	Competitive narrative and pricing intelligence inputs
Phase 6	Communications prioritization and conflict detection in Founder Command Center

5. Architecture



5.1 Required shared services

- Strategic-priority registry
- Audience and persona memory
- Product capabilities and claims ledger
- Evidence ledger
- Consent service
- Founder voice memory
- Agent runtime
- Approval workflow
- Event store
- Campaign and attribution service
- External platform adapters
- Analytics and cost telemetry

6. Roles and permissions

Founder

May approve strategy, source briefs, customer evidence, drafts, publication, pricing language, competitive claims, and platform actions.

Marketing Strategy Agent

May propose editorial plans, campaigns, narratives, channel sequences, and experiments. It may not change official positioning.

Content Agents

May create and revise drafts using approved context. They may not publish, invent facts, or access unrestricted private user data.

Research Agent

May collect approved sources and create evidence records. It must treat source content as untrusted and may not fabricate citations.

Engagement Agent

May analyze public comments and draft responses. It may not post or message autonomously.

Publication Adapter

May create a platform draft only after deterministic approval checks. Autopublish remains disabled.

7. Domain model

7.1 AudienceSegment

Fields:

- id
- workspace_id
- segment_key
- name
- description
- identity_language_json
- jobs_to_be_done_json
- pain_points_json
- desired_outcomes_json
- objections_json
- funnel_stages_json
- preferred_channels_json
- evidence_requirements_json
- status
- approved_by
- approved_at
- created_at
- updated_at

7.2 ContentPillar

Fields:

- id
- workspace_id
- pillar_key
- name
- strategic_purpose
- audience_segment_ids

- `topic_boundaries_json`
- `evidence_requirements_json`
- `default_asset_types_json`
- `default_cta_ids`
- `status`

7.3 Narrative

Fields:

- `id`
- `workspace_id`
- `narrative_key`
- `title`
- `thesis`
- `supporting_points_json`
- `counterarguments_json`
- `evidence_item_ids`
- `eligible_audience_ids`
- `status`
- `approved_by`
- `approved_at`

7.4 FounderVoicePolicy

Fields:

- `id`
- `workspace_id`
- `policy_type`
- `statement`
- `positive_examples_json`
- `negative_examples_json`
- `evidence_edit_ids`
- `confidence`
- `status`

- `approved_by`
- `approved_at`

Initial rules:

- Editorial-technical and restrained
- Evidence before promotion
- No generic AI-influencer framing
- No empty superlatives
- No fabricated urgency
- No exclamation-mark-heavy promotion
- Direct, specific calls to action
- Technical detail proportional to the audience
- Clear separation between observation, interpretation, and opinion

7.5 ChannelPolicy

Fields:

- `id`
- `workspace_id`
- `channel`
- `allowed_asset_types_json`
- `format_rules_json`
- `length_rules_json`
- `link_rules_json`
- `hashtag_rules_json`
- `publication_approval_required`
- `platform_draft_allowed`
- `autopublish_allowed`
- `status`

7.6 CallToAction

Fields:

- id
- workspace_id
- cta_key
- label
- destination_type
- destination_reference
- eligible_segments_json
- eligible_funnel_stages_json
- tracking_template
- active_from
- active_until
- status

7.7 ContentOpportunity

Fields:

- id
- workspace_id
- source_type
- source_entity_ids
- proposed_topic
- why_now
- business_objective
- audience_segment_ids
- content_pillar_id
- evidence_item_ids
- priority_score
- status
- created_at

Source types:

- Founder idea
- Product release
- Beta outcome
- User question
- Support pattern
- Research finding
- Competitive change
- Architecture decision
- QA result
- Event or webinar

7.8 ResearchBrief

Fields:

- id
- workspace_id
- content_opportunity_id
- research_question
- thesis_hypothesis
- source_plan_json
- required_evidence_json
- counterargument_plan_json
- status
- approved_by
- approved_at

7.9 SourceRecord

Fields:

- id
- workspace_id
- source_type
- title

- `author_or_publisher`
- `url_or_internal_reference`
- `published_at`
- `observed_at`
- `content_hash`
- `verification_status`
- `allowed_uses_json`
- `notes`

7.10 EvidenceMatrix

Fields:

- `id`
- `workspace_id`
- `research_brief_id`
- `claim_or_argument`
- `supporting_source_ids`
- `contradicting_source_ids`
- `confidence`
- `open_questions_json`
- `status`

7.11 SourceBrief

Fields:

- `id`
- `workspace_id`
- `content_opportunity_id`
- `research_brief_id`
- `title`
- `audience_segment_id`
- `funnel_stage`
- `content_pillar_id`
- `narrative_id`

- `problem`
- `insight`
- `proof`
- `implication`
- `primary_cta_id`
- `claims_manifest_json`
- `evidence_manifest_json`
- `prohibited_angles_json`
- `status`
- `approved_by`
- `approved_at`

7.12 Campaign

Fields:

- `id`
- `workspace_id`
- `campaign_key`
- `name`
- `objective`
- `audience_segment_ids`
- `offer_code`
- `start_at`
- `end_at`
- `channel_sequence_json`
- `asset_ids`
- `success_metrics_json`
- `budget_cents`
- `status`
- `approved_by`
- `approved_at`

7.13 ContentAsset

Fields:

- `id`
- `workspace_id`
- `source_brief_id`
- `campaign_id`
- `parent_asset_id`
- `asset_type`
- `channel`
- `title`
- `subtitle`
- `body`
- `structured_content_json`
- `claims_manifest_json`
- `evidence_manifest_json`
- `consent_manifest_json`
- `cta_id`
- `status`
- `version`
- `created_by_agent_run_id`
- `approval_id`
- `publication_record_id`
- `created_at`
- `updated_at`

Asset types:

- LinkedIn post
- LinkedIn carousel script
- Substack paper
- Substack newsletter
- Website page
- Landing-page section
- Email
- Email sequence

- Webinar outline
- Webinar script
- Release note
- Build log
- Case study
- Technical paper
- FAQ
- Lead magnet
- Comment response

7.14 EditorialCalendarItem

Fields:

- id
- workspace_id
- campaign_id
- content_asset_id
- event_type
- draft_due_at
- approval_due_at
- publish_at
- timezone
- dependency_ids
- status
- owner_id

7.15 PublicationRecord

Fields:

- id
- workspace_id
- content_asset_id
- channel
- external_id

- `external_url`
- `published_at`
- `published_by`
- `final_snapshot_json`
- `tracking_parameters_json`

7.16 EngagementRecord

Fields:

- `id`
- `workspace_id`
- `publication_record_id`
- `engagement_type`
- `external_actor_reference`
- `content_excerpt`
- `observed_at`
- `classification`
- `lead_signal`
- `response_draft_id`
- `status`

7.17 ContentPerformance

Fields:

- `id`
- `workspace_id`
- `publication_record_id`
- `measured_at`
- `impressions`
- `engagements`
- `comments`
- `clicks`
- `subscribers`
- `leads`

- `qualified_leads`
- `applications`
- `calls_booked`
- `enrollments`
- `revenue_cents`
- `attribution_method`
- `confidence`

7.18 ContentExperiment

Fields:

- `id`
- `workspace_id`
- `hypothesis`
- `variable_type`
- `variants_json`
- `audience_segment_id`
- `success_metric`
- `start_at`
- `end_at`
- `status`
- `result_json`
- `decision_record_id`

8. Agent specifications

8.1 Editorial Strategy Agent

Agent key: `fos.mcom.editorial_strategy`

Objective: produce a decision-oriented editorial plan aligned with current business priorities and available evidence.

Output schema:

```

{
  "period_start": "ISO-8601",
  "period_end": "ISO-8601",
  "business_objective": "string",
  "priority_topics": [
    {
      "topic": "string",
      "audience_segment_id": "uuid",
      "content_pillar_id": "uuid",
      "funnel_stage": "string",
      "why_now": "string",
      "recommended_source_asset": "string",
      "recommended_channels": ["string"],
      "primary_cta_id": "uuid",
      "required_evidence_ids": ["uuid"],
      "confidence": 0.0
    }
  ],
  "content_to_pause": ["string"],
  "founder_decisions_required": ["string"]
}

```

Guardrails:

- May not invent strategic priorities.
- Must identify insufficient evidence.
- Must not fill the calendar only because a slot is empty.
- Must distinguish demand generation, trust building, education, activation, and retention goals.

8.2 LinkedIn Content Agent

Agent key: `fos.mcom.linkedin`

Supported modes:

- Point-of-view post
- Educational post
- Build log
- Product demonstration
- Beta invitation
- Objection response
- Release post
- Carousel script
- Follow-up post

Output requirements:

- Hook or opening claim
- Core argument
- Supporting evidence
- Founder interpretation
- Clear CTA where appropriate
- Claims and evidence manifest
- Suggested first comment only when useful
- Alternative shorter version

Carousel output must include:

- Title slide
- One idea per slide
- Seven slides by default unless configured otherwise
- Visual hierarchy notes
- Diagram or data requirements
- Final CTA slide
- Accessibility text

8.3 Substack Research and Essay Agent

Agent key: `fos.mcom.substack`

Objective: produce rigorous, useful long-form work from a founder-approved research brief.

Output stages:

1. Research plan
2. Evidence matrix
3. Argument map
4. Outline
5. Section drafts
6. Counterargument section
7. Technical examples
8. Diagram briefs
9. Executive summary

10. Claims and source audit
11. Publication package
12. Derivative content plan

The agent must not create a final paper directly from a vague topic without an approved thesis and evidence plan.

8.4 Campaign Planning Agent

Agent key: `fos.mcom.campaign`

Produces:

- Campaign objective
- Audience
- Offer
- Sequence
- Asset dependencies
- Timing
- CTAs
- Measurement
- Risks
- Founder decisions

8.5 Repurposing Agent

Agent key: `fos.mcom.repurpose`

Rules:

- Derivatives inherit the approved claim set.
- A derivative may remove claims but may not add unsupported claims.
- Channel-native restructuring is required.
- Long-form nuance must not be converted into false certainty.
- Each derivative receives its own CTA and validation.

8.6 Claims, Consent, Pricing, and CTA Evaluator

Agent key: `fos.mcom.compliance`

Deterministic checks must verify:

- Claim approval and expiration
- Capability availability
- Current pricing
- Consent scope
- CTA activity and destination
- Channel policy
- Campaign dates
- Artificial scarcity
- Competitive evidence requirements

The model evaluator may flag semantic risk, but deterministic code makes the pass or block decision.

8.7 Founder Voice Evaluator

Agent key: `fos.mcom.founder_voice`

Inputs:

- Agent draft
- Founder final version
- Edit diff
- Approval reason
- Historical approved assets

Outputs:

- Proposed voice preference
- Supporting edit evidence
- Confidence
- Applicable asset types
- Possible overfitting risk

No voice rule becomes authoritative without founder approval.

8.8 Engagement Intelligence Agent

Agent key: `fos.mcom.engagement`

May:

- Cluster comments and replies
- Extract questions and objections
- Identify potential lead signals
- Draft founder responses
- Recommend follow-up content

May not:

- Post
- Message
- Infer sensitive attributes
- Treat engagement as purchase intent without evidence

9. Workflows

9.1 Content opportunity intake

1. Receive founder idea, release event, beta outcome, support theme, research finding, or strategic alert.
2. Create ContentOpportunity.
3. Verify privacy and public-use eligibility.
4. Score business impact, timeliness, evidence readiness, and founder effort.
5. Route high-value opportunities to Editorial Strategy Agent.

9.2 LinkedIn post workflow

1. Approve source brief.

2. Select post mode.
3. Generate draft.
4. Validate claims, evidence, CTA, channel policy, and voice.
5. Founder edits and approves.
6. Optionally create LinkedIn platform draft.
7. Founder publishes.
8. Record publication and performance.
9. Capture questions, leads, and follow-up opportunities.

9.3 LinkedIn carousel workflow

1. Approve source argument.
2. Generate seven-slide script.
3. Verify one idea per slide.
4. Generate design brief using the founder design system.
5. Approve script before visual generation.
6. Create final carousel artifact through the approved design workflow.
7. Verify accessibility and claim consistency.
8. Publish after founder approval.

9.4 Substack paper workflow

1. Approve research question and thesis.
2. Register sources.
3. Build evidence matrix.
4. Build argument map.
5. Approve outline.
6. Draft sections.
7. Run counterargument and factual review.
8. Run claims, product, pricing, consent, and voice review.
9. Founder revises and approves.
10. Create publication package.
11. Generate derivative LinkedIn, email, webinar, and website assets.

12. Track subscriber, lead, and enrollment outcomes.

9.5 Beta launch campaign workflow

Required sequence options:

- Problem-awareness post
- Founder thesis or why-now paper
- Product demonstration
- Beta invitation
- FAQ and objection response
- Webinar invitation
- Webinar follow-up
- Final call only when a factual deadline exists

Every campaign must specify one primary conversion event and measurement method.

9.6 Release communications workflow

1. Receive approved ReleaseCandidate and ReleaseReadinessReport.
2. Extract user-facing change.
3. Exclude internal-only or security-sensitive detail.
4. Generate release note, build log, LinkedIn draft, and customer update where appropriate.
5. Revalidate all affected claims.
6. Founder approves publication.

9.7 Engagement workflow

1. Import or record public engagement.
2. Classify question, objection, praise, lead signal, or noise.
3. Draft response where useful.
4. Route sensitive or consequential responses to founder.
5. Create content opportunity from repeated themes.

10. APIs

Strategy and configuration

- GET /api/fos/mcom/audiences
- POST /api/fos/mcom/audiences
- GET /api/fos/mcom/content-pillars
- POST /api/fos/mcom/content-pillars
- GET /api/fos/mcom/narratives
- POST /api/fos/mcom/narratives
- GET /api/fos/mcom/voice-policies
- POST /api/fos/mcom/voice-policies/:id/approve
- GET /api/fos/mcom/channel-policies
- GET /api/fos/mcom/ctas

Opportunities and research

- POST /api/fos/mcom/opportunities
- GET /api/fos/mcom/opportunities
- POST /api/fos/mcom/opportunities/:id/create-research-brief
- POST /api/fos/mcom/research-briefs/:id/approve
- POST /api/fos/mcom/research-briefs/:id/build-evidence-matrix
- GET /api/fos/mcom/source-briefs/:id

Campaigns and assets

- POST /api/fos/mcom/campaigns
- GET /api/fos/mcom/campaigns
- POST /api/fos/mcom/campaigns/:id/generate-plan
- POST /api/fos/mcom/assets/generate
- GET /api/fos/mcom/assets/:id
- PATCH /api/fos/mcom/assets/:id
- POST /api/fos/mcom/assets/:id/validate
- POST /api/fos/mcom/assets/:id/repurpose

- POST /api/fos/mcom/assets/:id/request-approval
- POST /api/fos/mcom/assets/:id/create-platform-draft

Calendar, publication, and performance

- GET /api/fos/mcom/calendar
- POST /api/fos/mcom/calendar-items
- POST /api/fos/mcom/publications
- POST /api/fos/mcom/performance/import
- GET /api/fos/mcom/dashboard
- GET /api/fos/mcom/attribution
- GET /api/fos/mcom/engagement
- POST /api/fos/mcom/engagement/:id/generate-response

11. User interface

11.1 Marketing Command Center

Shows:

- Current enrollment objective
- Campaigns
- Editorial priorities
- Drafts awaiting approval
- Research briefs
- Publication schedule
- Content opportunities
- Engagement requiring founder response
- Claim or consent blocks
- Performance and enrollment attribution

11.2 Editorial Calendar

Required views:

- Week
- Month
- Campaign sequence
- Channel
- Draft deadline
- Approval deadline
- Publication date
- Dependencies
- Status

11.3 LinkedIn Studio

Shows:

- Source brief
- Audience
- Funnel stage
- Post type
- Draft and alternatives
- Carousel slides where applicable
- Claims and evidence
- CTA
- Founder edit history
- Platform-draft action

11.4 Substack Paper Workspace

Tabs:

- Research question
- Sources
- Evidence matrix
- Argument map

- Outline
- Draft
- Counterarguments
- Diagrams
- Claims audit
- Publication package
- Derivatives
- Performance

11.5 Campaign Workspace

Shows:

- Objective
- Audience
- Offer
- Sequence
- Assets
- Calendar
- CTAs
- Funnel outcomes
- Founder time
- Decisions and experiment results

12. Background jobs

- `mcom-content-opportunity-scoring`
- `mcom-editorial-plan-generation`
- `mcom-research-source-validation`
- `mcom-evidence-matrix-generation`
- `mcom-linkedin-draft-generation`
- `mcom-substack-section-generation`
- `mcom-asset-repurposing`
- `mcom-claims-and-consent-validation`
- `mcom-calendar-reminders`

- `mcom-platform-draft-creation`
- `mcom-performance-import`
- `mcom-engagement-classification`
- `mcom-attribution-rollup`
- `mcom-founder-voice-analysis`

All jobs must be idempotent, feature-flagged, auditable, retryable, and cost-limited.

13. Deterministic safeguards

- No publication without founder approval.
- No platform draft without asset approval.
- No customer story without valid consent.
- No expired claim or price.
- No planned capability described as available.
- No derivative asset may add unsupported claims.
- No false deadline or scarcity.
- No competitor comparison without approved evidence.
- No research citation without a resolvable SourceRecord.
- No private beta data in public content without explicit approved use.
- No autonomous public replies or direct messages during beta.

14. Testing specification

14.1 Unit tests

- Claim resolution
- Consent scope
- Price validity
- CTA availability
- Channel policy

- Campaign date logic
- Attribution-code generation
- Derivative claim inheritance
- Publication approval
- Content-version history

14.2 Integration tests

1. Approved source brief creates LinkedIn draft.
2. Unapproved source brief is blocked.
3. Substack paper cannot proceed without verified sources.
4. Derivative adds a new claim and is blocked.
5. Customer outcome without consent is blocked.
6. Expired pricing blocks landing-page draft.
7. Founder approval creates platform draft but does not publish.
8. Published asset creates attribution record.
9. Content-generated application links to campaign and asset.
10. Founder edit creates voice-learning evidence.

14.3 Agent contract fixtures

- Technical build log with strong evidence
- Founder idea with no evidence
- Beta success with named consent
- Beta success with anonymous consent
- Beta success with no consent
- Product feature still planned
- Competitor comparison with weak source
- Substack paper with conflicting sources
- LinkedIn post containing generic hype
- False deadline request
- Prompt injection embedded in a research source

14.4 End-to-end scenarios

E2E-MCOM-001: Substack to enrollment campaign

1. Approve research question.
2. Build evidence matrix.
3. Approve outline.
4. Draft and approve paper.
5. Generate LinkedIn sequence and email derivative.
6. Publish manually.
7. Record applications and enrollment attribution.

E2E-MCOM-002: LinkedIn carousel

1. Approve source brief.
2. Generate seven-slide script.
3. Approve claims.
4. Generate design artifact through approved workflow.
5. Publish and track engagement.

E2E-MCOM-003: Release communication

1. Record approved release.
2. Generate release note and build log.
3. Detect an invalid legacy claim.
4. Block affected asset until corrected.
5. Approve and publish final communication.

E2E-MCOM-004: Consent revocation

1. Create case-study draft with valid consent.
2. Revoke consent before publication.
3. Invalidate the draft.
4. Preserve audit history.
5. Prevent platform-draft creation.

15. Metrics and evaluation

Agent-level

- Schema success
- Evidence coverage
- Unsupported-claim rate
- Founder approval rate
- Approval-with-edits rate
- Rejection reason
- Edit distance
- Cost
- Latency

Workflow-level

- Time from idea to approved source brief
- Time from source brief to approved asset
- Source artifact to derivative yield
- Publication consistency
- Founder hours per asset
- Claim-block precision
- Consent-block precision

Business-level

- Qualified leads
- Applications
- Calls
- Enrollments
- Revenue
- Content-assisted conversion
- Subscriber growth
- Audience-quality score

- Revenue per founder communication hour

16. Feature flags

- `fos_mcom_enabled`
- `fos_mcom_editorial_strategy_enabled`
- `fos_mcom_linkedin_enabled`
- `fos_mcom_carousel_enabled`
- `fos_mcom_substack_enabled`
- `fos_mcom_campaign_enabled`
- `fos_mcom_repurposing_enabled`
- `fos_mcom_engagement_enabled`
- `fos_mcom_platform_drafts_enabled`
- `fos_mcom_autopublish_enabled` - must remain false during beta

17. Work packages

WP-MCOM-0.1 - Strategy and configuration foundation

Deliverables:

- Audience segments
- Content pillars
- Narratives
- Founder voice policies
- Channel policies
- CTAs
- Configuration UI
- Tests

WP-MCOM-0.2 - Claims, consent, pricing, and publication gates

Deliverables:

- Validation service
- Publication approval
- Platform-draft boundary
- Audit trail
- Tests

WP-MCOM-1.1 - Beta launch campaign

Deliverables:

- Campaign model
- Campaign Planning Agent
- Launch sequence
- Calendar
- Attribution

WP-MCOM-1.2 - LinkedIn operating workflow

Deliverables:

- LinkedIn Content Agent
- Post types
- Carousel script schema
- LinkedIn Studio
- Platform-draft adapter

WP-MCOM-1.3 - Substack research and publication workflow

Deliverables:

- Research brief

- Source registry
- Evidence matrix
- Argument map
- Paper workspace
- Derivative generation

WP-MCOM-2.1 - Recurring editorial engine

Deliverables:

- Editorial Strategy Agent
- Weekly plan
- Content opportunity scoring
- Editorial calendar
- Founder workload limits

WP-MCOM-2.2 - Repurposing and engagement intelligence

Deliverables:

- Repurposing Agent
- Engagement import
- Question and objection clustering
- Response drafts

WP-MCOM-3.1 - Release, case study, and technical paper workflows

Deliverables:

- Release communication
- Case-study evidence package
- Architecture paper source brief
- Public changelog

WP-MCOM-4.1 - Attribution, experiments, and optimization

Deliverables:

- Performance import
- Lead and enrollment attribution
- Experiment registry
- Marketing dashboard
- Decision records

18. Deployment sequence

1. Migrate strategy and configuration entities.
2. Seed founder voice, audience, content pillars, channels, and CTAs.
3. Integrate claims, pricing, consent, and approvals.
4. Enable source briefs and LinkedIn drafts in shadow mode.
5. Enable Substack research workflow.
6. Enable beta launch campaign.
7. Enable platform-native draft creation after approval.
8. Enable editorial calendar and recurring plan.
9. Enable engagement intelligence.
10. Enable attribution and experiments.
11. Keep autopublish disabled until a separate governance decision.

19. Definition of done

MCOM is complete for beta when:

- Audience, pillar, narrative, voice, channel, CTA, and campaign records exist.
- A beta launch campaign can be generated and approved.

- LinkedIn posts and carousel scripts can be created from approved source briefs.
- Substack papers use research briefs, evidence matrices, outlines, and claims review.
- One approved source artifact can create multiple channel-native derivatives.
- Customer evidence is blocked without consent.
- Product, price, and competitive claims are validated.
- Platform drafts require approval.
- Autopublish is disabled.
- Publication and attribution records connect content to applications, enrollments, and revenue where evidence exists.
- Founder edits produce reviewable voice-learning proposals.
- Unit, integration, contract, security, and end-to-end tests pass.

20. Coding-agent execution instruction

Implement the Marketing and Communications Operating System according to this specification. Reuse the existing FOS identity, workspace isolation, memory, evidence, claims, consent, event, agent runtime, approval, telemetry, feature-flag, and deployment services.

Implement deterministic publication, claim, consent, pricing, CTA, and channel-policy gates before enabling content agents.

Begin with Phase 0 configuration and Phase 1 beta-launch workflows. Enable LinkedIn and Substack in shadow mode, then promote each workflow after contract tests pass.

Do not create a generic autonomous content bot. Preserve audience, source brief, evidence, campaign, channel, CTA, approval, publication, attribution, and founder-edit records.

Public publishing, public replies, direct messages, pricing changes, and customer-story use remain founder-controlled.